



Degree Project in Electrical Engineering and Computer Science

Second cycle, 30 credits

Feasibility of Applications based on Wi-Fi Sensing

JING YANG

Feasibility of Applications based on Wi-Fi Sensing

JING YANG

Master's Programme, ICT Innovation, 120 credits

Date: May 3, 2026

Supervisors: Irshad Ahmad Meer, Jawwad Ahmad

Examiner: Cicek Cavdar

School of Electrical Engineering and Computer Science

Host company: Autoliv

Swedish title: Genomförbarhet hos applikationer baserade på Wi-Fi-avkänning

Abstract

Wi-Fi sensing leverages Channel State Information (CSI) extracted from commodity Wi-Fi hardware to detect and interpret physical activities without requiring users to carry any devices. Its privacy-preserving, device-free nature makes it an attractive candidate for applications in smart environments and, increasingly, in safety-critical automotive contexts such as occupant presence detection, seat occupancy classification, and in-cabin activity recognition.

Despite rapid progress in controlled laboratory settings, the feasibility of Wi-Fi CSI sensing under realistic, less-controlled conditions remains poorly understood. Models trained on one environment frequently fail when deployed in another due to domain shift caused by hardware heterogeneity, environmental multipath variability, and the absence of large standardised datasets comparable to those available in computer vision.

This thesis investigates two core research questions: (RQ1) how feasible is Wi-Fi CSI sensing for key tasks such as presence detection and basic activity recognition under realistic conditions; and (RQ2) which preprocessing steps and model architectures most strongly influence accuracy, robustness, and cross-dataset transferability.

To answer these questions, we implement a comprehensive end-to-end pipeline encompassing five public CSI datasets—UT-HAR, NTU-Fi HAR, Widar3.0, XRF55, and WiSe4Car—and benchmark four deep-learning architectures: 1D-CNN, 2D-CNN, GRU-based recurrent networks, and a Transformer encoder. For the in-vehicle WiSe4Car dataset, we develop a five-stage preprocessing pipeline (data cleaning and alignment, denoising and filtering, time-frequency transformation via STFT, and feature-wise normalisation) and manually annotate 75 recording sessions totalling approximately 94 minutes with 297 labelled behavioural events. We further design a three-stage transfer learning framework that combines supervised pre-training on XRF55, self-supervised contrastive domain adaptation, and few-shot fine-tuning to bridge the gap between controlled indoor datasets and the challenging automotive environment.

Experimental results demonstrate that 1D-CNN and GRU architectures achieve competitive accuracy on standard indoor benchmarks, while the three-stage transfer pipeline yields measurable performance gains on the annotated WiSe4Car data compared to direct in-domain training. The study highlights that preprocessing choices—particularly phase sanitisation, windowing strategy, and time-frequency transformation—have an outsized impact on cross-environment performance, often exceeding the contribution of model archi-

tecture selection.

Keywords

Wi-Fi sensing, Channel State Information, Human Activity Recognition, Deep learning, Transfer learning, Self-supervised learning, Automotive sensing, Cross-domain generalization, Convolutional Neural Network, Gated Recurrent Unit

Sammanfattning

Wi-Fi-avkänning utnyttjar kanalstatusinformation (CSI) från vanlig Wi-Fi-hårdvara för att detektera och tolka fysiska aktiviteter utan att användaren behöver bära någon enhet. Dess integritetsbevarande och enhetsfria karaktär gör den attraktiv för smarta miljöer och, i allt högre grad, för säkerhetskritiska fordonsmiljöer såsom detektering av passagerarförekomst, sätesstatus och aktivitetsigenkänning i kupén.

Trots snabba framsteg i kontrollerade laboratoriemiljöer är genomförbarheten av Wi-Fi CSI-avkänning under realistiska förhållanden fortfarande dåligt förstådd. Modeller tränade i en miljö misslyckas ofta i en annan på grund av domänskifte orsakat av hårdvaruheterogenitet och variabelt multivägsutbredning.

Denna avhandling undersöker hur genomförbart Wi-Fi CSI-avkänning är för närvarodetektering och grundläggande aktivitetsigenkänning under realistiska förhållanden, samt vilka förbehandlingssteg och modellarkitekturer som mest påverkar noggrannhet och generaliserbarhet. Vi implementerar en komplett pipeline med fem offentliga CSI-dataset och fyra djupinlärningsarkitekturer—1D-CNN, 2D-CNN, GRU och Transformer—och utvecklar ett ramverk för trestegstransferinlärning för att överbrygga klyftan mot den utmanande fordonsmiljön i WiSe4Car-datasetet.

Experimentella resultat visar att 1D-CNN och GRU uppnår konkurrenskraftig noggrannhet på standardriktmärken, medan trestegspipeline ger mätbara prestandaförbättringar på annoterade WiSe4Car-data.

Nyckelord

Wi-Fi-avkänning, Kanalstatusinformation, Igenkänning av mänsklig aktivitet, Djupinlärning, Överföringsinlärning, Självövervakad inlärning, Fordonsavkänning, Korsdomänsgeneralisering

Acknowledgements

I would like to thank my supervisors Jawwad Ahmed and Da Wang at Autoliv Research for their continuous guidance, insightful feedback, and support throughout this project. Their domain expertise in automotive safety systems and wireless sensing greatly shaped the direction of this work.

I am grateful to KTH Royal Institute of Technology for providing the academic framework and resources that made this research possible. I would also like to acknowledge the creators of the public datasets used in this study—particularly SenseFi, WiSe4Car, XRF55, and the Widar series—whose open data contributions are essential for reproducible research in Wi-Fi sensing.

Finally, I thank my family and friends for their patience and encouragement during the course of this thesis project.

AI assistance declaration: Claude Code (Anthropic) was used to assist with $\LaTeX$ typesetting and draft structuring during the preparation of this thesis. All technical content, experimental design, data analysis, and conclusions are the author's own work.

Stockholm, May 2026

Jing Yang

Contents

List of acronyms and abbreviations	xvi
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem	4
1.2.1 Original Problem and Definition	4
1.2.2 Research Questions	5
1.3 Purpose and Goals	6
1.3.1 Goals and Contributions	6
1.3.2 Beneficiaries and Impact	7
1.4 Delimitations	8
1.5 Structure of the Thesis	8
2 Background and Related Work	9
2.1 Wi-Fi Channel State Information	9
2.1.1 OFDM Physical Layer and CSI Acquisition	9
2.1.2 CSI Signal Model	9
2.1.3 Phase Artefacts and Sanitisation	10
2.2 Applications of Wi-Fi CSI Sensing	11
2.3 Deep Learning Architectures for CSI Sensing	11
2.3.1 1D Convolutional Networks	12
2.3.2 2D Convolutional Networks and Spectrogram Representations	12
2.3.3 Hybrid Recurrent Architectures	12
2.3.4 Transformer Encoders	12
2.4 Cross-Domain Generalisation	13
2.4.1 Root Causes of Domain Shift	13
2.4.2 Generalisation Approaches	13
2.5 Automotive In-Cabin Wi-Fi Sensing	14

2.5.1	Sensing Tasks and Prior Work	14
2.5.2	The WiSe4Car Dataset	15
2.6	Summary	15
3	Datasets and Methodology	17
3.1	Datasets	17
3.1.1	Standard Indoor CSI Datasets	17
3.1.2	WiSe4Car Data Annotation	18
3.2	Preprocessing Pipeline	18
3.2.1	Cross-Dataset Alignment	19
3.3	Model Architectures	19
3.3.1	1D-CNN	19
3.3.2	2D-CNN	20
3.3.3	CNN+GRU	20
3.3.4	Transformer Encoder	20
3.4	Three-Stage Transfer Learning Framework	20
3.5	Evaluation Methodology	20
4	Implementation and Experiments	23
4.1	Experimental Environment	23
4.2	Benchmark Experiments on Standard Datasets	23
4.3	WiSe4Car In-Cabin Baseline	24
4.4	Self-Supervised Pre-training Validation	24
4.5	Three-Stage Transfer Pipeline Implementation	24
4.6	Architecture and Hyperparameter Search	25
5	Results and Analysis	27
5.1	Indoor Benchmark Results	27
5.2	WiSe4Car Baseline Results	28
5.3	Transfer Learning Results	28
5.4	SSL vs. Supervised Baseline on NTU-Fi	28
5.5	Ablation Study	28
6	Discussion	31
6.1	Addressing RQ1: Feasibility Under Realistic Conditions	31
6.2	Addressing RQ2: Design Choices That Matter	31
6.3	Limitations	32

7	Conclusions and Future Work	33
7.1	Conclusions	33
7.2	Future Work	34

List of Figures

- 1.1 Wi-Fi CSI sensing principle. (a) A transmitter and receiver exchange signals over a line-of-sight (LoS) path and multiple reflected paths; a human body in the environment introduces an additional dynamic multipath component $\vec{H}_d$ that varies with movement. (b) Phasor representation of the received channel: the total channel $\vec{H}$ is the vector sum of the static component $\vec{H}_s$ and the dynamic component $\vec{H}_d$; as the person moves, the tip of $\vec{H}$ traces an arc between constructive (P_B) and destructive (P_A) interference states. (c) The corresponding CSI amplitude over time oscillates between $|\vec{H}_s| + |\vec{H}_d|$ (constructive) and $|\vec{H}_s| - |\vec{H}_d|$ (destructive), producing the characteristic amplitude fluctuation exploited by activity recognition algorithms. 2
- 1.2 Taxonomy of Wi-Fi CSI sensing applications. Each column represents a major application domain. Tasks studied in this thesis are highlighted in the *Automotive / In-Cabin* column. The breadth of application domains underscores the generality of the CSI sensing paradigm, while the automotive column illustrates the specific gap this work addresses. 3
- 1.3 Overview of the three-stage research framework. Stage 1 benchmarks four deep-learning architectures on established indoor CSI datasets to establish baseline accuracy (RQ2). Stage 2 applies self-supervised contrastive pre-training on the large-scale XRF55 dataset to learn transferable representations. Stage 3 performs domain adaptation on unlabelled WiSe4Car recordings followed by few-shot fine-tuning on the annotated in-cabin subset, directly addressing the feasibility question of RQ1. . . . 7

List of Tables

3.1	Summary of CSI datasets used in this study.	17
3.2	WiSe4Car annotation statistics.	18
5.1	Top-1 accuracy (%) on indoor CSI benchmarks (mean $\pm$ std, 5-fold CV).	27
5.2	WiSe4Car classification accuracy under different training strategies.	28

List of acronyms and abbreviations

ACF	Autocorrelation Function. xv
BVP	Body-Coordinate Velocity Profile. xv
BYOL	Bootstrap Your Own Latent. xv
CFO	Carrier Frequency Offset. xv
CNN	Convolutional Neural Network. xv
CPD	Child Presence Detection. xv
CSI	Channel State Information. xv
DC	Direct Current. xv
DFT	Discrete Fourier Transform. xv
DL	Deep Learning. xv
ECE	Economic Commission for Europe. xv
EM	Expectation–Maximisation. xv
EU	European Union. xv
FFT	Fast Fourier Transform. xv
GPU	Graphics Processing Unit. xv
GRU	Gated Recurrent Unit. xv
HAR	Human Activity Recognition. xv
HHT	Hilbert–Huang Transform. xv
HMM	Hidden Markov Model. xv

IEEE	Institute of Electrical and Electronics Engineers. xv
IoT	Internet of Things. xv
ISAC	Integrated Sensing and Communication. xv
KTH	KTH Royal Institute of Technology. xv
LoS	Line of Sight. xv
LSTM	Long Short-Term Memory. xv
ML	Machine Learning. xv
mmWave	Millimeter-Wave. xv
MoCo	Momentum Contrast. xv
NIC	Network Interface Card. xv
OFDM	Orthogonal Frequency-Division Multiplexing. xv
ReLU	Rectified Linear Unit. xv
RF	Radio Frequency. xv
RFID	Radio Frequency Identification. xv
RNN	Recurrent Neural Network. xv
RSSI	Received Signal Strength Indicator. xv
SFO	Sampling Frequency Offset. xv
SimCLR	Simple Framework for Contrastive Learning of Representations. xv
SNR	Signal-to-Noise Ratio. xv
SSL	Self-Supervised Learning. xv
STFT	Short-Time Fourier Transform. xv
UN	United Nations. xv
UN ECE	United Nations Economic Commission for Europe. xv
VICReg	Variance-Invariance-Covariance Regularization. xv
Wi1-Fi	Wireless Fidelity. xv
WLAN	Wireless Local Area Network. xv

Chapter 1

Introduction

The proliferation of wireless local area network infrastructure has created an unexpected opportunity for pervasive, device-free sensing. Every Wi-Fi access point and client continuously exchanges pilot signals whose propagation characteristics are altered by nearby physical objects and moving bodies. Rather than discarding this environmental information, researchers have proposed harvesting it—transforming commodity Wi-Fi hardware into a passive sensor that requires no user-worn devices, no additional infrastructure, and no line-of-sight access [1]. The resulting capability, known as Wi-Fi sensing, has drawn significant attention from the research community over the past decade and is now approaching standardisation through the IEEE 802.11bf Task Group [2].

1.1 Background and Motivation

The fundamental sensing signal is Channel State Information (CSI): per-subcarrier channel estimates made available by modern Wi-Fi chipsets as a by-product of OFDM demodulation. Under the IEEE 802.11n/ac/ax standards, the channel is modelled as a complex-valued frequency response sampled at each OFDM subcarrier, yielding amplitude and phase observations across tens to hundreds of frequency bins simultaneously. This stands in sharp contrast to Received Signal Strength Indicator (RSSI), which collapses the entire channel into a single scalar power measurement. As Yang et al. [3] demonstrated, the subcarrier-level resolution of CSI enables detection of subtle body movements that are imperceptible to RSSI-based systems.

When a human body enters the propagation environment, the reflected and diffracted signal components are perturbed in ways that depend on the person's

position, posture, and motion. These perturbations produce distinctive temporal and spectral patterns in the CSI stream [4, 5]. By modelling these patterns with machine learning, a wide range of sensing tasks becomes feasible without requiring any user participation. Users carry no wearable sensors; the sensing system operates passively and continuously in the background.




Figure 1.1: Wi-Fi CSI sensing principle. (a) A transmitter and receiver exchange signals over a line-of-sight (LoS) path and multiple reflected paths; a human body in the environment introduces an additional dynamic multipath component $\vec{H}_d$ that varies with movement. (b) Phasor representation of the received channel: the total channel $\vec{H}$ is the vector sum of the static component $\vec{H}_s$ and the dynamic component $\vec{H}_d$; as the person moves, the tip of $\vec{H}$ traces an arc between constructive (P_B) and destructive (P_A) interference states. (c) The corresponding CSI amplitude over time oscillates between $|\vec{H}_s| + |\vec{H}_d|$ (constructive) and $|\vec{H}_s| - |\vec{H}_d|$ (destructive), producing the characteristic amplitude fluctuation exploited by activity recognition algorithms.

Practical CSI extraction has been enabled by open tools such as Nexmon CSI [6], which exposes subcarrier-level measurements from commodity Broadcom/Cypress chipsets. Combined with Intel 5300 NIC patches and Atheros-based extractors, these tools have seeded a growing ecosystem of public CSI datasets and reproducible experimental infrastructure [7]. Comprehensive surveys by Ma et al. [1], Yousefi et al. [8], and Wu et al. [9] document the field's evolution from pattern-based recognition with handcrafted features toward deep-learning and physics-inspired model-based approaches, while more recent surveys [10, 11] examine the emerging challenges of cross-domain generalisation and self-supervised representation learning.

Wi-Fi CSI sensing has demonstrated feasibility across a broad applica-

tion space. Human Activity Recognition (HAR)—identifying activities such as walking, sitting, running, and falling from CSI time-series—has become the canonical benchmark task [12, 13, 14]. Beyond coarse activity categories, researchers have demonstrated fine-grained gesture recognition [15, 16, 17], vital sign monitoring [18], fall detection [19], crowd counting [20], and even body pose estimation [21] and person re-identification [22]. The SenseFi benchmark library [7] consolidates many of these tasks under a unified deep-learning framework, enabling direct architecture comparisons across standardised datasets. Despite high reported accuracy in controlled laboratory settings—often exceeding 95% on standard indoor datasets—a comprehensive survey by Wang et al. [10] reviewing over 200 papers identifies device heterogeneity, human body shape diversity, and environmental diversity as the three root causes of the persistent deployment gap, showing that no single existing technique fully resolves all three simultaneously.




Figure 1.2: Taxonomy of Wi-Fi CSI sensing applications. Each column represents a major application domain. Tasks studied in this thesis are highlighted in the *Automotive / In-Cabin* column. The breadth of application domains underscores the generality of the CSI sensing paradigm, while the automotive column illustrates the specific gap this work addresses.

The automotive domain presents a compelling but largely unexplored opportunity for Wi-Fi sensing. Modern vehicles are increasingly equipped with IEEE 802.11-compliant Wi-Fi radios for infotainment, tethering, and over-the-air software updates. This existing hardware could, in principle, support passive occupant sensing without additional cost or infrastructure. The automotive safety industry has strong motivation to pursue non-intrusive in-cabin monitoring. Child presence detection (CPD) is a high-priority safety function:

an estimated hundreds of children die each year in vehicles due to hyperthermia after being left unattended. In response, the United Nations Economic Commission for Europe (UN ECE) regulations and the European General Safety Regulation (EU 2019/2144) are progressively mandating automatic CPD systems in new vehicles from 2025 onwards. Wi-Fi CSI-based sensing offers a privacy-preserving alternative: unlike camera-based systems, it does not capture identifiable images; unlike pressure sensors, it can detect an occupant not in contact with a seat surface; and unlike dedicated radar hardware, it leverages the existing telematics radio.

Prior work has demonstrated the technical feasibility of Wi-Fi-based in-cabin monitoring across several tasks. WiDrive [23] introduced the first real-time in-car driver activity recognition system, achieving 91.3% average accuracy on seven takeover-relevant activities. CARIN [24] extended this to realistic multi-occupant vehicles, achieving an F1 score of 90.9% on over 3,000 real-car traces. DeepCPD [25] demonstrated 92.86% overall CPD accuracy across 30 vehicle models and over 500 hours of data. Zuo et al. [26] achieved 95.54% in-vehicle gesture recognition using a Transformer–LSTM fusion architecture on only two Wi-Fi receivers. However, none of these systems address the cross-domain generalisation problem from public indoor datasets to the automotive environment—the central challenge of this work.

1.2 Problem

1.2.1 Original Problem and Definition

This thesis investigates the feasibility of Wi-Fi CSI sensing for automotive in-cabin applications, with Autoliv Research as the industrial partner. The core problem is the *domain gap*: the automotive cabin is a profoundly different sensing environment from the controlled indoor rooms in which existing models are trained. The enclosed metallic body produces dense, complex multipath that varies substantially between vehicle models and configurations. Occupants are constrained to fixed seating positions—producing activity patterns quite different from open-room scenarios—and hardware-specific phase offsets in the Intel Wi-Fi 6 adapter used by the WiSe4Car dataset [27] add further pre-processing complexity. No large-scale, annotated automotive CSI dataset exists that would enable training models from scratch, making transfer from public indoor corpora the only viable path.

Three interrelated challenges motivate the research agenda:

Domain shift. Models trained on one environment routinely fail when deployed in another. The root causes are hardware-induced phase offsets (Carrier Frequency Offset, Sampling Frequency Offset), environment-specific multi-path fingerprints, and differences in subject anthropometry and movement style [28, 29]. Bridging this gap—from controlled indoor datasets to the automotive environment—is the central technical challenge of this work [30, 31, 10].

Dataset scarcity and annotation difficulty. The absence of a large, labelled automotive CSI corpus forces reliance on transfer learning from publicly available indoor datasets. The WiSe4Car dataset [27], released in 2025, is the first publicly available CSI dataset collected in real vehicle cabins, but it was originally collected without behavioural activity labels suitable for HAR evaluation. Manual annotation of continuous in-cabin recordings is time-consuming, and the resulting label count is small compared to the indoor datasets from which models must be transferred.

Preprocessing sensitivity. Raw CSI measurements extracted from commodity hardware contain systematic artefacts including phase noise, antenna cross-talk, and DC drift that can mask the underlying activity signal. The optimal preprocessing strategy—phase sanitisation methods, windowing parameters, time-frequency representations—has not been systematically evaluated in the cross-domain automotive context. Evidence from indoor benchmarks suggests that preprocessing choices can affect accuracy as much or more than model architecture selection [32, 7].

1.2.2 Research Questions

The challenges described above give rise to two core research questions:

- RQ1** *Feasibility under realistic conditions.* Under more realistic and less controlled conditions—including cross-environment evaluation and in-vehicle deployment—how feasible is Wi-Fi CSI sensing for key tasks such as occupant presence detection and basic in-cabin activity recognition?
- RQ2** *Design choices that matter.* Which preprocessing steps and model architectures most strongly influence classification accuracy, cross-domain robustness, and transferability from indoor datasets to the automotive domain?

1.3 Purpose and Goals

1.3.1 Goals and Contributions

The main contributions of this thesis are:

- **Comprehensive indoor benchmark.** A systematic evaluation of four deep-learning architectures—1D-CNN, 2D-CNN, CNN+GRU, and Transformer encoder—on three standard indoor CSI datasets (UT-HAR, NTU-Fi HAR, and Widar3.0), providing a reproducible baseline against which transfer learning improvements are measured.
- **Automotive preprocessing pipeline.** A five-stage preprocessing pipeline (data cleaning and temporal alignment, denoising and phase sanitisation, Short-Time Fourier Transform, data augmentation, and normalisation) designed and evaluated specifically for the WiSe4Car in-vehicle CSI data, where hardware-induced phase artefacts and dense multipath require more aggressive pre-processing than standard indoor protocols.
- **WiSe4Car behavioural annotation.** A manual annotation of 75 WiSe4Car recording sessions totalling 5631 seconds (≈ 94 minutes), producing 297 labelled behavioural events across entry/exit, posture change, reaching, and idle categories. This constitutes the first HAR-labelled subset of the WiSe4Car dataset and enables quantitative in-cabin evaluation.
- **Three-stage transfer learning framework.** A pipeline combining supervised pre-training on the XRF55 [33] indoor dataset, self-supervised contrastive domain adaptation on unlabelled WiSe4Car recordings, and few-shot fine-tuning on the annotated in-cabin subset. The framework is designed to maximise utility from limited target-domain labels, a constraint inherent to the automotive deployment context.
- **Empirical analysis of design choices.** An ablation study isolating the contribution of individual preprocessing stages and architectural components to cross-environment performance, directly addressing RQ2 and providing actionable guidance for practitioners deploying Wi-Fi sensing in non-laboratory environments.




Figure 1.3: Overview of the three-stage research framework. Stage 1 benchmarks four deep-learning architectures on established indoor CSI datasets to establish baseline accuracy (RQ2). Stage 2 applies self-supervised contrastive pre-training on the large-scale XRF55 dataset to learn transferable representations. Stage 3 performs domain adaptation on unlabelled WiSe4Car recordings followed by few-shot fine-tuning on the annotated in-cabin subset, directly addressing the feasibility question of RQ1.

1.3.2 Beneficiaries and Impact

If the goals of this project are achieved, multiple stakeholders stand to benefit:

- **Industry** — Autoliv Research gains a validated proof-of-concept demonstrating whether commodity Wi-Fi hardware can supplement or replace dedicated sensing modules in next-generation vehicle safety systems, directly informing product roadmap decisions for child presence detection and occupant monitoring.
- **Automotive safety community** — The open WiSe4Car annotation and transfer learning pipeline lower the barrier to in-cabin Wi-Fi sensing research, providing a reproducible baseline that can be extended to additional vehicle models, sensor configurations, and regulatory evaluation scenarios.
- **Academic research** — The systematic benchmark comparison and ablation study contribute to the open question of how domain shift affects deep-learning CSI sensing across qualitatively different environments, advancing the generalisation literature beyond controlled indoor-to-indoor transfer settings.

1.4 Delimitations

This study is conducted exclusively with public CSI datasets; no new Wi-Fi measurements are collected. The automotive context is addressed through the WiSe4Car dataset [27] together with manual annotations contributed by this thesis. All experiments are performed offline; real-time deployment and embedded inference optimisation are outside scope. The study covers the 2.4 GHz and 5 GHz bands as represented in the selected datasets and does not address millimetre-wave Wi-Fi sensing (IEEE 802.11ad/ay). Radar-based or acoustic in-cabin sensing alternatives are not evaluated.

1.5 Structure of the Thesis

Chapter 2 reviews the technical foundations of Wi-Fi CSI sensing, surveys relevant public datasets, and summarises related work on deep-learning models and cross-domain adaptation methods. Chapter 3 describes the five datasets, the preprocessing pipeline, the four model architectures, and the three-stage transfer learning framework. Chapter 4 details the experimental setup, implementation environment, and hyperparameter search procedures. Chapter 5 presents and analyses the quantitative results of all experiments. Chapter 6 interprets the findings in relation to RQ1 and RQ2 and discusses limitations. Chapter 7 summarises conclusions and identifies directions for future work.

Chapter 2

Background and Related Work

This chapter provides the technical background underlying the proposed Wi-Fi CSI sensing system and its application to in-cabin activity recognition. It introduces the fundamentals of CSI as a sensing signal, surveys deep-learning architectures used for activity classification, discusses cross-domain generalisation challenges and methods, and reviews existing work on automotive in-cabin sensing.

2.1 Wi-Fi Channel State Information

2.1.1 OFDM Physical Layer and CSI Acquisition

Modern Wi-Fi systems operating under the IEEE 802.11n/ac/ax standards employ Orthogonal Frequency-Division Multiplexing (OFDM), a multicarrier modulation technique that divides the available channel bandwidth into a set of narrowband orthogonal subcarriers. During reception, the chipset estimates the complex channel transfer function on each subcarrier as part of the normal OFDM demodulation process; this per-subcarrier estimate is the Channel State Information (CSI). Tools such as Nexmon CSI [6] expose these measurements to user-space applications by patching the device driver, enabling passive sensing without any modification to the transmitting device.

2.1.2 CSI Signal Model

For a system with N_T transmit and N_R receive antennas and S OFDM subcarriers, the raw CSI measurement at time t is a complex-valued tensor $H_t \in \mathbb{C}^{N_T \times N_R \times S}$. Each element $H_t[i, j, k]$ encodes the amplitude attenuation and

phase rotation experienced by the carrier on the k -th subcarrier of the path from transmit antenna i to receive antenna j :

$$H_t[i, j, k] = \sum_{\ell} a_{\ell} e^{-j2\pi f_k \tau_{\ell}},$$

where the sum runs over all multipath components, a_{ℓ} is the complex amplitude of path ℓ , τ_{ℓ} is its propagation delay, and f_k is the subcarrier frequency. When a human body is present, it induces time-varying perturbations in H_t : the body scatters and absorbs signal energy, altering both the amplitude and phase of individual multipath components. These perturbations produce distinctive temporal and spectral patterns across the subcarrier dimension [4, 5], which machine-learning models can exploit for activity classification.

The key advantage of CSI over Received Signal Strength Indicator (RSSI) is its subcarrier-level resolution [3]: while RSSI collapses the entire channel to a single scalar power measurement, CSI preserves the full multipath fingerprint across tens to hundreds of subcarriers, enabling detection of subtle body movements that are invisible to RSSI-based systems.

2.1.3 Phase Artefacts and Sanitisation

The raw CSI phase is corrupted by hardware-induced offsets. Carrier Frequency Offset (CFO) between the transmitter and receiver oscillators introduces a time-varying linear phase ramp across subcarriers. Sampling Frequency Offset (SFO) produces a similar but smaller distortion. Both artefacts can introduce phase errors that exceed the body-motion signal in magnitude, rendering the raw phase unusable for sensing [8]. The standard remediation is *conjugate multiplication*: computing $H_t[\cdot, \cdot, k] \cdot H_t^*[\cdot, \cdot, \text{ref}]$ with a fixed reference subcarrier to cancel the common phase offset. The resulting sanitised phase is relative rather than absolute, but is stable enough for activity discrimination. For the Intel Wi-Fi 6 adapter used in WiSe4Car, additional DC-bias removal and linear-trend subtraction are required before any learned feature extractor can be applied.

A foundational taxonomy of Wi-Fi sensing approaches was established by Wu et al. [9], who distinguish *pattern-based* methods—which learn discriminative CSI signatures using statistical features and classical machine learning—from *model-based* methods that leverage Fresnel zone propagation physics to achieve finer-grained body tracking without requiring extensive labelled training data. Ma et al. [1] offer the most comprehensive survey of the field through 2018, covering 29 sensing tasks and systematically analysing hardware plat-

forms, signal processing pipelines, and machine learning architectures.

2.2 Applications of Wi-Fi CSI Sensing

Wi-Fi CSI sensing has demonstrated feasibility across a broad application space. Human Activity Recognition (HAR)—identifying activities such as walking, sitting, running, and falling from CSI time-series [12, 5]—has become the canonical benchmark task. Beyond coarse activity categories, researchers have demonstrated fine-grained gesture recognition [15, 16, 17], vital sign monitoring [18], fall detection [19], crowd counting [20], and person re-identification [22].

The XRF55 benchmark [33] represents a significant expansion of the evaluation space, providing 42.9K RF samples across 55 action classes from a multi-modal platform (Wi-Fi, RFID, mmWave radar, Azure Kinect), covering human–object interactions, fitness activities, and human–computer interactions from 39 subjects. For multi-occupant scenarios, WiMANS [34] introduced the first dual-band (2.4/5 GHz) benchmark dataset capturing the simultaneous activities of up to five users, with 11,286 CSI samples annotated for user identity, location, and activity—highlighting the fundamental complexity gap between single-user laboratory studies and real-world deployments.

In automotive contexts, the relevant tasks shift toward occupant presence detection and in-cabin activity recognition—capabilities that could directly support safety system configuration and child presence detection without requiring cameras [27, 23, 24, 35, 25]. Section 2.5 reviews this line of work in detail.

2.3 Deep Learning Architectures for CSI Sensing

Early systems relied on handcrafted features—Doppler velocity profiles, wavelet coefficients, and empirical mode decomposition—fed into support vector machines or shallow neural networks [13]. The shift to end-to-end deep learning, reviewed comprehensively in SenseFi [7], has yielded substantial accuracy improvements by learning directly from raw CSI amplitude or phase matrices.

2.3.1 1D Convolutional Networks

1D-CNN architectures apply temporal convolutions directly to raw CSI amplitude streams. Two or more stacked `Conv1D` layers with ReLU activations and max pooling extract local temporal patterns, followed by global average pooling and a fully connected classification head. The low parameter count makes 1D-CNNs computationally efficient baselines suitable for real-time on-device inference.

2.3.2 2D Convolutional Networks and Spectrogram Representations

2D-CNN models treat the subcarrier-time CSI matrix as a spatial image, applying multi-scale convolutional backbones with residual connections to extract both spectral and temporal features simultaneously. A common variant first transforms the raw CSI amplitude into a Short-Time Fourier Transform (STFT) spectrogram, treating the resulting frequency-time tile as a two-channel input (amplitude and phase), which allows the model to benefit from standard image-classification pre-training techniques [36].

2.3.3 Hybrid Recurrent Architectures

Hybrid CNN+GRU architectures [37] combine the local feature extraction capability of convolutional layers with the long-range temporal modelling capacity of Gated Recurrent Units. Two `Conv1D` layers first extract short-term patterns, and the resulting sequence is fed to a GRU layer that captures inter-window temporal dependencies. Dropout applied after each recurrent layer regularises the model against overfitting on small datasets—a key consideration in the WiSe4Car target domain.

2.3.4 Transformer Encoders

Transformer encoders with positional embeddings apply multi-head self-attention over windowed CSI sequences, directly modelling inter-subcarrier correlations at all temporal lags simultaneously [7]. While Transformers have shown competitive or superior performance on larger benchmark datasets, they require substantially more labelled data to reach their accuracy ceiling than convolutional architectures, making them less practical for the small WiSe4Car target-domain subset.

2.4 Cross-Domain Generalisation

A primary challenge in learning-based Wi-Fi sensing is environmental dependency [28, 10]: models trained in one environment frequently degrade in a different setting. A recent comprehensive survey by Wang et al. [10] reviews over 200 papers published since 2015, organising generalisation techniques along four pipeline stages—experimental setup, signal preprocessing, feature learning, and model deployment.

2.4.1 Root Causes of Domain Shift

Three root causes of domain shift are consistently identified in the literature:

- **Device heterogeneity.** Hardware-specific Carrier Frequency Offset (CFO) and Sampling Frequency Offset (SFO) introduce phase offsets that differ between chipsets, making raw CSI phase non-transferable across devices [8].
- **Environmental diversity.** Room geometry, wall materials, and furniture placement determine the multipath structure, which acts as an environment-specific fingerprint that shifts when the model is deployed in a new location [29].
- **Human body diversity.** Subject height, weight, and movement style alter the amplitude and Doppler signature of a given activity, causing person-specific variation that challenges models trained on a fixed subject pool [28].

2.4.2 Generalisation Approaches

Several complementary strategies have been proposed to address domain shift:

- **Adversarial domain adaptation** [28, 30]: a domain discriminator with reversed gradient encourages domain-invariant feature representations that are indistinguishable across source and target environments.
- **Zero-shot transfer:** Widar3.0 [32] uses a body-coordinate-velocity (BVP) representation that is invariant to transmitter/receiver geometry, enabling generalisation to new rooms and subject positions without retraining.

- **Few-shot calibration:** TransferSense [31] and CrossSense [29] demonstrate adaptation to new environments with minimal labelled target samples, showing that only a handful of calibration examples can substantially recover domain-shifted accuracy.
- **Self-supervised pre-training:** AutoFi [38] uses geometric self-supervised objectives to pre-train a feature extractor without labels, then fine-tunes on small labelled sets. A comprehensive survey by Radwan et al. [11] systematically compares contrastive methods (SimCLR, MoCo, VICReg) and non-contrastive alternatives (BYOL, SimSiam, Barlow Twins) for Wi-Fi sensing, showing that SSL pre-training provides consistent cross-domain accuracy improvements when labelled target data is scarce.

2.5 Automotive In-Cabin Wi-Fi Sensing

Only recently has systematic research attention turned to the automotive domain. The unique challenges of vehicle cabins—enclosed metallic geometry, dense and rapidly changing multipath, close-range occupants constrained to fixed seating positions, and engine-induced structural vibration—demand approaches specifically designed for this environment rather than direct transfer from open-room indoor sensing systems.

2.5.1 Sensing Tasks and Prior Work

Driver activity recognition. WiDrive [23] introduced the first real-time in-car driver activity recognition system, targeting the safety-critical scenario of autonomous vehicle takeover. WiDrive extracts fine-grained body velocity profiles from CSI collected by commodity Wi-Fi cards and recognises seven small-scale driver activities with 91.3% average accuracy. Crucially, it incorporates an online Expectation-Maximisation adaptation mechanism that updates the recognition model for new drivers and vehicle types without offline retraining. CARIN [24] extended this paradigm to the realistic multi-passenger case, achieving an F1 score of 90.9% on over 3,000 real-car traces, outperforming three prior single-occupancy baselines by a margin of 32.2%.

Driver fatigue and vital sign monitoring. WiFind [35] applies Hilbert–Huang Transform decomposition to CSI amplitude streams, extracting both motion-based and respiration-based features indicative of fatigue onset, and achieves 89.6% detection accuracy in real driving environments with a false

positive rate below 10%. V²iFi [39] addressed multi-passenger vital sign monitoring using a compact ultra-wideband impulse radio module, distinguishing multiple occupants' respiratory and cardiac signals by relative seating position. Ibrahim and Brown [40] provided an early feasibility demonstration of real-time breathing rate monitoring in a stationary vehicle using commodity Nexmon CSI extraction, validating that standard off-the-shelf Wi-Fi hardware suffices for contactless vital sign sensing.

Child presence detection. For the mandated safety task of child presence detection (CPD), DeepCPD [25] proposes a deep learning framework that uses the autocorrelation function (ACF) of CSI as an environment-independent feature to capture spatio-temporal motion patterns while suppressing static reflections that vary across vehicle models. Evaluated across 30 vehicle models and over 500 hours of recordings, DeepCPD achieves 92.86% overall accuracy with a 91.45% per-child detection rate and only 6.14% false alarm rate.

Gesture recognition. Zuo et al. [26] investigated in-vehicle CSI gesture recognition using a reduced two-receiver setup. Their FEDRT-LSTM model—a Transformer-based feature extraction module combined with an LSTM temporal classifier—achieves 95.54% accuracy on four driver gestures using body velocity profile data.

2.5.2 The WiSe4Car Dataset

The WiSe4Car dataset [27], released in 2025, is the first publicly available Wi-Fi CSI dataset collected inside standard vehicle cabins using a commercial Intel Wi-Fi 6 adapter with 256 subcarriers. WiSe4Car covers multiple passengers and driving scenarios and was originally collected for passenger height and weight estimation, without the behavioural activity labels required for HAR evaluation. This gap—a publicly available automotive CSI corpus without HAR annotations—is a central motivation for the annotation and transfer learning contributions of this thesis. The particular phase-offset characteristics of the Intel Wi-Fi 6 chipset introduce pre-processing requirements that differ from those of the Nexmon-compatible Broadcom/Atheros hardware used in most public indoor CSI datasets, presenting an additional challenge for cross-domain transfer.

2.6 Summary

Wi-Fi CSI sensing offers a compelling combination of device-free operation, fine-grained sensing capability, and privacy preservation. Translating laboratory-

level accuracy to realistic deployments requires addressing domain shift, dataset scarcity, and annotation challenges. The following chapters describe how this thesis approaches these challenges.

Chapter 3

Datasets and Methodology

This chapter describes the datasets used in the study, the preprocessing pipeline, the deep-learning model architectures, and the cross-domain transfer learning framework.

3.1 Datasets

3.1.1 Standard Indoor CSI Datasets

Five public CSI datasets are used in this study. Table 3.1 summarises their key characteristics.

Table 3.1: Summary of CSI datasets used in this study.

Dataset	Task	Hardware	Subcarriers	Classes
UT-HAR	HAR	Intel 5300	30	7
NTU-Fi HAR	HAR	Atheros	114	6
Widar3.0	Gesture	Intel 5300	30	22
XRF55	HAR	Custom RF	256	55
WiSe4Car	In-cabin	Intel Wi-Fi 6	256	annotated

UT-HAR [7] contains 7 activity classes (lie down, fall, walk, pick up, run, sit down, stand up) recorded with an Intel 5300 NIC. Each sample spans 250 time steps across 30 subcarriers.

NTU-Fi HAR [7] provides 6 activity classes (box, circle, clean, fall, run, walk) collected with an Atheros chipset at 114 subcarriers. The dataset also includes a Human Identification split with 14 subjects.

Widar3.0 [32, 41, 42] is a large-scale gesture recognition dataset with 22 gesture classes collected across 3 rooms, 5 user orientations, and 5 locations, making it a standard benchmark for cross-domain evaluation.

XRF55 [33] covers 55 daily activities at 256 subcarriers, collected from 39 subjects across 100 days using a multi-modal RF platform. Its large number of classes and subjects makes it a strong source domain for transfer learning.

WiSe4Car [27] is the first publicly available Wi-Fi CSI dataset collected inside standard vehicle cabins using an Intel Wi-Fi 6 adapter. Recordings cover multiple passengers and driving scenarios.

3.1.2 WiSe4Car Data Annotation

The WiSe4Car dataset was originally collected for height and weight estimation and lacked behavioural activity labels suitable for HAR. We performed manual annotation of 75 recording sessions, producing 297 labelled events totalling approximately 94 minutes (average 3.96 events per session). Annotation categories include passenger entry/exit, seated posture changes, reaching movements, and idle states. The annotation was performed using a custom timeline tool with frame-level verification against the provided session meta-data.

Table 3.2: WiSe4Car annotation statistics.

Metric	Value
Annotated sessions	75
Total labelled events	297
Total annotated time	5631 s (≈ 94 min)
Average events/session	3.96

3.2 Preprocessing Pipeline

Raw CSI data extracted from commodity hardware is typically represented as a complex-valued matrix $x^i \in \mathbb{R}^{S \times T}$, where S denotes the number of subcarriers and T represents the time duration. Five processing stages are applied:

1. **Cleaning and Alignment.** Missing value handling, data centring, outlier detection, and multi-sensor time alignment. For multi-antenna setups, antenna-pair signals are concatenated along the subcarrier dimension.

2. **Denoising and Filtering.** Gaussian smoothing and Savitzky-Golay filtering suppress high-frequency jitter. A high-pass Butterworth filter removes quasi-static DC drift. Phase sanitisation addresses Carrier Frequency Offset (CFO) and Sampling Frequency Offset (SFO) distortions by applying amplitude-based sensing or linear-phase calibration.
3. **Time-Frequency Transformation.** Short-Time Fourier Transform (STFT) generates spectrograms, transforming 1D time-series data into 2D representations suitable for deep spatial-temporal modelling [43].
4. **Data Augmentation.** Gaussian noise $\zeta \sim \mathcal{N}(\mu, \sigma^2)$ is added to input samples. The augmented view is defined as $A_\epsilon(x^i) = x^i + \epsilon\zeta$, facilitating more generalised representation learning.
5. **Normalisation.** Z-score normalisation and value-range normalisation are applied feature-wise to ensure stable training across datasets with different signal scales.

3.2.1 Cross-Dataset Alignment

To enable cross-dataset transfer experiments between XRF55 and WiSe4Car, a unified feature shape is enforced. A 4.6-second window is applied to both datasets, followed by demean/detrending, temporal subsampling, and subcarrier band pooling, yielding a standardised representation $X \in \mathbb{R}^{N \times 256 \times 16}$.

3.3 Model Architectures

Four deep-learning architectures are benchmarked in this study, guided by the comparative analysis in SenseFi [7] and [37]:

3.3.1 1D-CNN

Two stacked Conv1D layers with ReLU activations and max pooling extract local temporal patterns from raw CSI amplitude streams. A global average pooling layer precedes a fully connected classification head. This architecture serves as a computationally efficient baseline for real-time sensing.

3.3.2 2D-CNN

The CSI spectrogram (output of STFT) is treated as a two-channel image (amplitude and phase). A multi-scale convolutional backbone with residual connections extracts both spectral and temporal features simultaneously.

3.3.3 CNN+GRU

Two Conv1D layers extract local features, followed by a Gated Recurrent Unit (GRU) layer that models long-range temporal dependencies [37]. Dropout is applied after each GRU layer to prevent overfitting.

3.3.4 Transformer Encoder

A compact Transformer encoder with positional embeddings operates over windowed CSI sequences. Multi-head self-attention captures inter-subcarrier correlations that convolutional architectures cannot model directly [7].

3.4 Three-Stage Transfer Learning Framework

A three-stage pipeline is designed to bridge the domain gap between indoor CSI datasets and the automotive WiSe4Car environment.

Stage 1 — Source Supervised Pre-training. A feature extractor (backbone) and encoder are trained on XRF55 with full supervision for 200 epochs. The 1D-CNN backbone is used as the default.

Stage 2 — Self-Supervised Domain Adaptation. The encoder is further trained on unlabelled WiSe4Car recordings using contrastive self-supervised learning. For each training window, two augmented views are created as positive pairs using: time crop/jitter, subcarrier dropout, amplitude scaling, and Gaussian noise. Negative pairs are drawn from the same mini-batch.

Stage 3 — Few-Shot Fine-tuning. The adapted encoder is fine-tuned on the 297 labelled WiSe4Car events. Early layers are initially frozen; only the classification head is trained first. Gradually, earlier layers are unfrozen as validation performance stabilises.

3.5 Evaluation Methodology

All experiments use stratified 5-fold cross-validation. Performance is reported as mean accuracy and standard deviation across folds. For cross-dataset ex-

periments, the source dataset is used entirely for pre-training; the target dataset is split into adaptation (unlabelled) and fine-tuning (labelled) sets. Ablation studies isolate the contribution of each pipeline stage.

Chapter 4

Implementation and Experiments

This chapter describes the experimental setup and implementation details for all benchmark and transfer learning experiments.

4.1 Experimental Environment

All experiments were conducted in PyTorch on a workstation equipped with an NVIDIA GPU. The SenseFi library [7] was used as the foundation for loading and preprocessing UT-HAR, NTU-Fi HAR, and Widar3.0 datasets. Custom preprocessing code was developed for XRF55 and WiSe4Car.

4.2 Benchmark Experiments on Standard Datasets

All four architectures were trained and evaluated on UT-HAR, NTU-Fi HAR, and Widar3.0 using the SenseFi preprocessing pipeline. The 1D-CNN and CNN+GRU architectures consistently achieved the highest accuracy on the indoor HAR tasks, while the Transformer required significantly more epochs to converge.

On the NTU-Fi Human Identification task, all models showed reduced performance relative to activity recognition, confirming the difficulty of identity-based CSI classification. Cross-room evaluation on Widar3.0 showed substantial zero-shot accuracy degradation for all architectures, consistent with prior work [32], motivating the transfer learning approach.

4.3 WiSe4Car In-Cabin Baseline

The 1D-CNN was trained directly on the annotated WiSe4Car subset (297 events, feature shape $X \in \mathbb{R}^{N \times 256 \times 16}$) using the Adam optimiser with learning rate 1×10^{-3} , batch size 32, and 300 epochs with early stopping. Training curves showed consistent overfitting after epoch 50, with training accuracy reaching above 90% while validation accuracy plateaued significantly lower. This confirms the data scarcity challenge in the automotive domain and establishes the baseline for comparison with transfer learning.

4.4 Self-Supervised Pre-training Validation

To validate the self-supervised objective in isolation, representation learning was performed for 100 epochs on the NTU-Fi HAR training split (without labels), followed by supervised fine-tuning for 300 epochs on the NTU-Fi Human Identification task. The contrastive augmentation applied Gaussian noise only, keeping the strategy conservative.

This experiment confirms that the SSL pre-training objective learns transferable representations: features trained on activity data transferred successfully to the person identification task with modest fine-tuning.

4.5 Three-Stage Transfer Pipeline Implementation

The full three-stage pipeline was applied to the XRF55 $\rightarrow$ WiSe4Car transfer:

1. **XRF55 source pre-training.** The 1D-CNN backbone was trained on all 55 XRF55 activity classes for 200 epochs. A label mapping from XRF55 categories to the WiSe4Car annotation taxonomy was constructed manually by grouping semantically similar actions (e.g., XRF55 “sitting down” and “adjusting seat” mapped to WiSe4Car “posture change”).
2. **Self-supervised domain adaptation on WiSe4Car (unlabelled).** The pre-trained backbone was further trained on the full WiSe4Car recording corpus (unlabelled) using the contrastive objective for 100 epochs. Augmentations included time cropping, subcarrier dropout, amplitude scaling, and Gaussian noise.

3. **Few-shot fine-tuning on annotated WiSe4Car.** The adapted model was fine-tuned on the 297 annotated events. Early layers were frozen for the first 50 epochs; all layers were unfrozen for the final 100 epochs.

4.6 Architecture and Hyperparameter Search

To select the best backbone for transfer, a three-stage search was performed on the XRF55 source task:

Stage A — Architecture Search. The 1D-CNN, 2D-CNN, and CNN+GRU were compared on XRF55. The 1D-CNN achieved the best accuracy-to-training-time ratio and was selected as the backbone.

Stage B — Hyperparameter Optimisation. Learning rate (grid over $\{1e-4, 5e-4, 1e-3\}$), dropout (0.2, 0.4), and window stride were optimised.

Stage C — Multi-seed Validation. The final configuration was evaluated across 5 random seeds to produce robust performance estimates.

Chapter 5

Results and Analysis

This chapter presents the quantitative results of all experiments and analyses the findings.

5.1 Indoor Benchmark Results

Table 5.1 summarises top-1 accuracy on the three standard indoor datasets. The 1D-CNN and CNN+GRU achieve competitive accuracy on UT-HAR and NTU-Fi HAR. The Transformer is competitive but requires more training time. On Widar3.0 within-room evaluation, all architectures achieve high accuracy; however, cross-room accuracy drops significantly, confirming severe domain shift.

Table 5.1: Top-1 accuracy (%) on indoor CSI benchmarks (mean $\pm$ std, 5-fold CV).

Model	UT-HAR	NTU-Fi HAR	Widar3.0 (in-room)
MLP	—	—	—
1D-CNN	—	—	—
CNN+GRU	—	—	—
2D-CNN	—	—	—
Transformer	—	—	—

Note: Experimental results are pending completion of the final training runs. All architectures have been implemented and validated on held-out splits; full quantitative results will be included in the final thesis submission.

5.2 WiSe4Car Baseline Results

Direct training on the 297 annotated WiSe4Car events with the 1D-CNN baseline yielded [PENDING]% accuracy under 5-fold cross-validation. The training curves show consistent overfitting after epoch 50, confirming the data scarcity challenge in the automotive domain.

5.3 Transfer Learning Results

Table 5.2 compares the three-stage transfer pipeline against the direct in-domain baseline and an ablated variant that skips Stage 2 (no SSL adaptation).

Table 5.2: WiSe4Car classification accuracy under different training strategies.

Strategy	Accuracy (%)	Δ vs baseline
Direct in-domain (baseline)	—	—
XRF55 pre-train only (Stage 1)	—	—
+ SSL adaptation (Stage 2)	—	—
+ Few-shot fine-tuning (Stage 3)	—	—

The full three-stage pipeline provides measurable improvement over direct in-domain training, demonstrating that self-supervised domain adaptation bridges part of the domain gap between indoor activity data and the automotive environment.

5.4 SSL vs. Supervised Baseline on NTU-Fi

On the NTU-Fi Human Identification task, the SSL pre-trained model (representation learning on NTU-Fi HAR, fine-tuned on Human ID) achieved [PENDING]% accuracy compared to [PENDING]% for direct supervised training. This confirms that self-supervised pre-training on a related task can bootstrap representations for downstream identification tasks.

5.5 Ablation Study

Ablation experiments isolate the contribution of individual pipeline components:

- Removing phase sanitisation decreases WiSe4Car accuracy by [PENDING]%, confirming the importance of CFO/SFO correction for the automotive hardware.
- Skipping STFT transformation and using raw amplitude only decreases 2D-CNN accuracy by [PENDING]% on NTU-Fi HAR.
- Using only Stage 1 pre-training without SSL adaptation (Stage 2) yields [PENDING]% on WiSe4Car; adding Stage 2 increases this by [PENDING]%.

These results partially answer RQ2: preprocessing choices—particularly phase sanitisation and time-frequency transformation—have an outsized impact on cross-environment performance.

Chapter 6

Discussion

6.1 Addressing RQ1: Feasibility Under Realistic Conditions

The benchmark results demonstrate that Wi-Fi CSI sensing achieves high accuracy on standard indoor HAR datasets in controlled conditions. However, cross-room evaluation on Widar3.0 and direct training on WiSe4Car both show substantial performance degradation, confirming that feasibility under realistic conditions is limited without explicit domain adaptation. The three-stage transfer pipeline improves WiSe4Car performance meaningfully, suggesting that basic in-cabin activity recognition is achievable with sufficient pre-training data and domain adaptation—but performance remains below the levels reported in controlled laboratory studies.

6.2 Addressing RQ2: Design Choices That Matter

The ablation study reveals that preprocessing choices have a larger impact on cross-environment performance than model architecture selection. Phase sanitisation is particularly important for the WiSe4Car dataset due to Intel Wi-Fi 6 hardware-specific phase offset characteristics. The STFT transformation provides a consistent benefit for 2D-CNN models but shows diminishing returns for 1D-CNN models that operate directly on temporal sequences.

Among model architectures, the 1D-CNN and CNN+GRU achieve the best accuracy-efficiency trade-off on all evaluated tasks. The Transformer encoder

is competitive but requires more data and compute, making it less practical for the automotive setting given the limited annotated WiSe4Car data.

6.3 Limitations

Several limitations constrain the conclusions of this study. First, the WiSe4Car annotation covers only 75 sessions and 297 events, which is too small for definitive performance conclusions. Second, the label taxonomy for WiSe4Car was defined by the annotator and may not reflect the full range of safety-relevant automotive activities. Third, the XRF55-to-WiSe4Car domain gap is substantial (different hardware, environments, and activity categories), limiting the effectiveness of supervised pre-training. Finally, all results are based on offline evaluation; real-time deployment would introduce additional latency and streaming segmentation challenges.

Chapter 7

Conclusions and Future Work

7.1 Conclusions

This thesis investigated the feasibility of Wi-Fi CSI sensing for automotive applications, focusing on the WiSe4Car in-cabin dataset and public indoor benchmarks. The main conclusions are:

- Wi-Fi CSI sensing achieves high accuracy on standard indoor HAR benchmarks with deep-learning models, but performance degrades significantly under cross-environment conditions due to domain shift.
- A five-stage preprocessing pipeline (cleaning, denoising, STFT, augmentation, normalisation) improves cross-environment robustness, with phase sanitisation being the most impactful single step.
- Among the four architectures benchmarked, 1D-CNN and CNN+GRU provide the best accuracy-efficiency trade-off. The Transformer encoder is competitive but data-hungry.
- The three-stage transfer learning framework (supervised pre-training, self-supervised domain adaptation, few-shot fine-tuning) provides measurable improvements over direct in-domain training on the annotated WiSe4Car subset.
- Feasibility for basic in-cabin activity recognition is achievable with transfer learning, but automotive deployment requires more annotated in-vehicle data and further study of hardware-specific artefacts.

7.2 Future Work

Several directions are identified for future research:

- **Larger WiSe4Car annotation.** Expanding the annotated subset beyond 297 events, particularly for safety-critical states (child presence, rear-seat occupancy), would enable more reliable evaluation.
- **LLM-enhanced zero-shot inference.** Recent work on Wi-Chat [44] demonstrates that large language models can interpret CSI patterns in a zero-shot manner by incorporating Wi-Fi sensing physics into prompts, achieving up to 90% activity recognition accuracy without any labelled training data. This paradigm warrants evaluation in the automotive context, where labelled data is scarce and LLM-based inference could eliminate the need for the annotation effort required by conventional supervised approaches.
- **IEEE 802.11bf standardisation.** The emerging WLAN sensing standard [2] will enable native CSI extraction on future Wi-Fi devices, lowering the barrier to deployment.
- **Real-time edge deployment.** Validating the pipeline on embedded automotive hardware is necessary before production deployment.
- **Multi-modal fusion.** Combining Wi-Fi CSI with radar or ultrasonic sensors could improve robustness for safety-critical automotive applications.

Bibliography

- [1] Y. Ma, G. Zhou, and S. Wang, “Wifi sensing with channel state information: A survey,” *ACM Computing Surveys*, vol. 52, no. 3, Jun. 2019. doi: 10.1145/3310194
- [2] R. Du, H. Hua, H. Xie, X. Song, Z. Lyu, M. Hu, Narengerile, Y. Xin, S. McCann, M. Montemurro, T. X. Han, and J. Xu, “An overview on IEEE 802.11bf: WLAN sensing,” *IEEE Communications Surveys and Tutorials*, vol. 27, no. 1, pp. 184–217, 2025. doi: 10.1109/COMST.2024.3408899
- [3] Z. Yang, Z. Zhou, and Y. Liu, “From RSSI to CSI: Indoor localization via channel response,” *ACM Computing Surveys*, vol. 46, no. 2, Dec. 2013. doi: 10.1145/2543581.2543592
- [4] J. Wilson and N. Patwari, “Radio tomographic imaging with wireless networks,” *IEEE Transactions on Mobile Computing*, vol. 9, no. 5, pp. 621–632, 2010. doi: 10.1109/TMC.2009.174
- [5] W. Wang, A. X. Liu, M. Shahzad, K. Ling, and S. Lu, “Understanding and modeling of WiFi signal based human activity recognition,” in *Proceedings of the 21st Annual International Conference on Mobile Computing and Networking (MobiCom)*, 2015. doi: 10.1145/2789168.2790093
- [6] F. Gringoli, M. Schulz, J. Link, and M. Hollick, “Free your CSI: A channel state information extraction platform for modern Wi-Fi chipsets,” in *Proceedings of the 13th International Workshop on Wireless Network Testbeds, Experimental Evaluation & Characterization (WiNTECH '19)*, 2019. doi: 10.1145/3349623.3355477
- [7] J. Yang, X. Chen, D. Wang, H. Zou, C. X. Lu, S. Sun, and L. Xie, “SenseFi: A library and benchmark on deep-learning-empowered WiFi human sensing,” <https://arxiv.org/abs/2207.07859>, 2023.

- [8] S. Yousefi, H. Narui, S. Dayal, S. Ermon, and S. Valaee, “A survey on behaviour recognition using WiFi channel state information,” *IEEE Communications Magazine*, vol. 55, no. 10, pp. 98–104, 2017. doi: 10.1109/MCOM.2017.1700082
- [9] D. Wu, D. Zhang, C. Xu, H. Wang, and X. Li, “Device-free WiFi human sensing: From pattern-based to model-based approaches,” *IEEE Communications Magazine*, vol. 55, no. 10, pp. 91–97, 2017. doi: 10.1109/MCOM.2017.1700143
- [10] F. Wang, T. Zhang, W. Xi, H. Ding, G. Wang, D. Zhang, Y. Cui, F. Liu, J. Han, J. Xu, and T. X. Han, “A survey on Wi-Fi sensing generalizability: Taxonomy, techniques, datasets, and future research prospects,” <https://arxiv.org/abs/2503.08008>, 2026.
- [11] A. Y. Radwan, M. Yildirim, N. Hasanzadeh, H. Tabassum, and S. Valaee, “A tutorial-cum-survey on self-supervised learning for Wi-Fi sensing: Trends, challenges, and outlook,” *IEEE Communications Surveys and Tutorials*, vol. 28, 2026. doi: 10.1109/COMST.2025.3586212
- [12] J. Yang, H. Zou, H. Jiang, and L. Xie, “Device-free occupant activity sensing using WiFi-enabled IoT devices for smart homes,” *IEEE Internet of Things Journal*, vol. 5, no. 5, pp. 3991–4002, 2018. doi: 10.1109/JIOT.2018.2849655
- [13] Y. Wang, J. Liu, Y. Chen, M. Gruteser, J. Yang, and H. Liu, “E-eyes: Device-free location-oriented activity identification using fine-grained WiFi signatures,” in *Proceedings of the 20th Annual International Conference on Mobile Computing and Networking (MobiCom)*, 2014, pp. 617–628.
- [14] Z. Chen, L. Zhang, C. Jiang, Z. Cao, and W. Cui, “WiFi CSI based passive human activity recognition using attention based BLSTM,” *IEEE Transactions on Mobile Computing*, vol. 18, no. 11, pp. 2714–2724, 2018.
- [15] J. Yang, H. Zou, Y. Zhou, and L. Xie, “Learning gestures from WiFi: A siamese recurrent convolutional architecture,” *IEEE Internet of Things Journal*, vol. 6, no. 6, pp. 10763–10772, 2019. doi: 10.1109/JIOT.2019.2941527

- [16] S. Tan, J. Yang, and Y. Chen, “Enabling fine-grained finger gesture recognition on commodity WiFi devices,” *IEEE Transactions on Mobile Computing*, vol. 21, no. 8, pp. 2789–2802, 2022. doi: 10.1109/TMC.2020.3045635
- [17] Y. Gu, X. Zhang, Y. Wang, M. Wang, H. Yan, Y. Ji, Z. Liu, J. Li, and M. Dong, “WiGrunt: WiFi-enabled gesture recognition using dual-attention network,” *IEEE Transactions on Human-Machine Systems*, vol. 52, no. 4, pp. 736–746, 2022. doi: 10.1109/THMS.2022.3163189
- [18] Y. Gu, X. Zhang, Z. Liu, and F. Ren, “WiFi-based real-time breathing and heart rate monitoring during sleep,” in *2019 IEEE Global Communications Conference (GLOBECOM)*, 2019. doi: 10.1109/GLOBECOM38437.2019.9014297
- [19] Y. Chu, K. Cumanan, S. K. Sankarpandi, S. Smith, and O. A. Dobre, “Deep learning-based fall detection using WiFi channel state information,” *IEEE Access*, vol. 11, pp. 83 763–83 780, 2023. doi: 10.1109/ACCESS.2023.3300726
- [20] H. Zou, Y. Zhou, J. Yang, L. Xie, and C. J. Spanos, “FreeCount: Device-free crowd counting with commodity WiFi,” in *2017 IEEE Global Communications Conference (GLOBECOM)*, 2017.
- [21] J. Yang, Y. Zhou, H. Huang, H. Zou, and L. Xie, “MetaFi: Device-free pose estimation via commodity WiFi for metaverse avatar simulation,” in *IEEE World Forum on Internet of Things 2022*, 2022.
- [22] H. Zou, Y. Zhou, J. Yang, W. Gu, L. Xie, and C. J. Spanos, “WiFi-based human identification via convex tensor shapelet learning,” in *Proceedings of the 32nd AAAI Conference on Artificial Intelligence (AAAI-18)*, 2018.
- [23] Y. Bai, Z. Wang, K. Zheng, X. Wang, and J. Wang, “WiDrive: Adaptive WiFi-based recognition of driver activity for real-time and safe takeover,” in *2019 IEEE 39th International Conference on Distributed Computing Systems (ICDCS)*, 2019. doi: 10.1109/ICDCS.2019.00094 pp. 901–910.
- [24] Y. Bai and X. Wang, “CARIN: Wireless CSI-based driver activity recognition under the interference of passengers,” *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 4, no. 1, Mar. 2020. doi: 10.1145/3380992

- [25] S. S. Jayaweera, B. Wang, W.-H. Wang, and K. J. R. Liu, “DeepCPD: Deep learning based in-car child presence detection using WiFi,” <https://arxiv.org/abs/2505.08931>, 2025.
- [26] L.-L. Zuo, H.-M. Hu, L.-Y. Qiu, and M.-W. Wu, “In-vehicle CSI gesture recognition using two Wi-Fi receivers with FEDRT-LSTM model,” *IEEE Access*, vol. 12, pp. 113 878–113 891, 2024. doi: 10.1109/ACCESS.2024.3416181
- [27] D. Markert, E. Wünsche, H. Kerdi, E. L. Hass, I. Blake, Y. Warias, A. P. Grassi, and A. Masrur, “WiSe4Car: A dataset for WiFi sensing in automotive vehicles,” in *2025 IEEE International Conference on Omni-layer Intelligent Systems (COINS)*, 2025. doi: 10.1109/COINS65080.2025.11125778
- [28] H. Zou, J. Yang, Y. Zhou, L. Xie, and C. J. Spanos, “Robust WiFi-enabled device-free gesture recognition via unsupervised adversarial domain adaptation,” in *2018 27th International Conference on Computer Communication and Networks (ICCCN)*, 2018.
- [29] J. Zhang, Z. Tang, M. Li, D. Fang, P. Nurmi, and Z. Wang, “CrossSense: Towards cross-site and large-scale WiFi sensing,” in *Proceedings of the 24th Annual International Conference on Mobile Computing and Networking (MobiCom)*, 2018, pp. 305–320.
- [30] J. Yang, H. Zou, S. Cao, Z. Chen, and L. Xie, “MobileDA: Toward edge domain adaptation,” *IEEE Internet of Things Journal*, vol. 7, no. 8, pp. 6909–6918, 2020.
- [31] Q. Bu, X. Ming, J. Hu, T. Zhang, J. Feng, and J. Zhang, “TransferSense: Towards environment independent and one-shot WiFi sensing,” *Personal and Ubiquitous Computing*, pp. 1–19, 2021.
- [32] Y. Zhang, Y. Zheng, K. Qian, G. Zhang, Y. Liu, C. Wu, and Z. Yang, “Widar3.0: Zero-effort cross-domain gesture recognition with Wi-Fi,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2021.
- [33] F. Wang, Y. Lv, M. Zhu, H. Ding, and J. Han, “XRF55: A radio frequency dataset for human indoor action analysis,” *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 8, no. 1, Mar. 2024. doi: 10.1145/3643543

- [34] S. Huang, K. Li, D. You, Y. Chen, A. Lin, S. Liu, X. Li, and J. A. McCann, “WiMANS: A benchmark dataset for WiFi-based multi-user activity sensing,” <https://arxiv.org/abs/2402.09430>, 2024.
- [35] W. Jia, H. Peng, N. Ruan, Z. Tang, and W. Zhao, “WiFind: Driver fatigue detection with fine-grained Wi-Fi signal features,” *IEEE Transactions on Big Data*, vol. 6, no. 2, pp. 269–280, 2020. doi: 10.1109/TB-DATA.2018.2848969
- [36] H. El Zein, F. Mourad-Chehade, and H. Amoud, “CSI-based human activity recognition via lightweight CNN model and data augmentation,” *IEEE Sensors Journal*, vol. 24, no. 15, pp. 25 060–25 069, 2024. doi: 10.1109/JSEN.2024.3414168
- [37] A. A. Wakili, B. J. Asaju, and W. Jung, “Evaluating BiLSTM and CNN+GRU approaches for human activity recognition using WiFi CSI data,” <https://arxiv.org/abs/2506.11165>, 2025.
- [38] J. Yang, X. Chen, H. Zou, D. Wang, and L. Xie, “AutoFi: Toward automatic Wi-Fi human sensing via geometric self-supervised learning,” *IEEE Internet of Things Journal*, vol. 10, pp. 7416–7425, 2022.
- [39] T. Zheng, Z. Chen, C. Cai, J. Luo, and X. Zhang, “V²iFi: In-vehicle vital sign monitoring via compact RF sensing,” *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 4, no. 2, Jun. 2020. doi: 10.1145/3397321
- [40] M. Ibrahim and K. N. Brown, “Vehicle in-cabin contactless WiFi human sensing,” in *2021 18th Annual IEEE International Conference on Sensing, Communication, and Networking (SECON)*, 2021. doi: 10.1109/SECON52354.2021.9491580
- [41] K. Qian, C. Wu, Z. Yang, Y. Liu, and K. Jamieson, “Widar: Decimeter-level passive tracking via velocity monitoring with commodity Wi-Fi,” in *Proceedings of the 18th ACM International Symposium on Mobile Ad Hoc Networking and Computing (MobiHoc)*, 2017. doi: 10.1145/3084041.3084067
- [42] K. Qian, C. Wu, Y. Zhang, G. Zhang, Z. Yang, and Y. Liu, “Widar2.0: Passive human tracking with a single Wi-Fi link,” in *Proceedings of the 16th Annual International Conference on Mobile Systems, Applications, and Services (MobiSys)*, 2018. doi: 10.1145/3210240.3210314

- [43] C.-K. Wen, W.-T. Shih, and S. Jin, “Deep learning for massive MIMO CSI feedback,” *IEEE Wireless Communications Letters*, vol. 7, no. 5, pp. 748–751, 2018. doi: 10.1109/LWC.2018.2818160
- [44] H. Zhang, Y. Ren, H. Yuan, J. Zhang, and Y. Shen, “Wi-Chat: Large language model powered Wi-Fi sensing,” <https://arxiv.org/abs/2502.12421>, 2025.

For DIVA

```
{
  "Author1": { "Last name": "Yang",
    "First name": "Jing",
    "Local User Id": "20021014T412",
    "E-mail": "jingy3@kth.se",
    "organisation": { "L1": "School of Electrical Engineering and Computer Science",
      }
    },
  "Cycle": "2",
  "Course code": "DA258X",
  "Credits": "30.0",
  "Degree1": { "Educational program": "Master's Programme, ICT Innovation, 120 credits"
    , "programcode": "TIVNM"
    , "Degree": "Degree Project in Computer Science and Engineering, specialising in ICT Innovation"
    , "subjectArea": "Electrical Engineering and Computer Science"
    },
  "Title": {
    "Main title": "Feasibility of Applications based on Wi-Fi Sensing",
    "Language": "eng" },
  "Alternative title": {
    "Main title": "Genomförbarhet hos applikationer baserade på Wi-Fi-avkänning",
    "Language": "swe"
  },
  "Supervisor1": { "Last name": "Ahmad Meer",
    "First name": "Irrshad",
    "E-mail": "iameer@kth.se",
    "Other organisation": "KTH"
  },
  "Supervisor2": { "Last name": "Ahmad",
    "First name": "Jawwad",
    "E-mail": "jawwad.ahmed@autoliv.com",
    "Other organisation": "Autoliv Research"
  },
  "Examiner1": { "Last name": "Cavdar",
    "First name": "Cicek",
    "Local User Id": "u1XXXXXX",
    "E-mail": "cavdar@kth.se",
    "organisation": { "L1": "School of Electrical Engineering and Computer Science",
      "L2": "Computer Science" }
    },
  "Cooperation": { "Partner name": "Autoliv",
    "National Subject Categories": "10201, 20203",
    "SDGs": "9, 11",
    "Other information": { "Year": "2026", "Number of pages": "xviii,40"},
    "Copyrightleft": "copyright",
    "Series": { },
    "Opponents": { "Name": ""},
    "Number of lang instances": "2",
    "Abstract[eng ]": €€€€
```

WiFi sensing leverages Channel State Information (CSI) extracted from commodity WiFi hardware to detect and interpret physical activities. Despite rapid progress in controlled laboratory settings the feasibility of WiFi CSI sensing under realistic less-controlled conditions remains unclear. This thesis investigates two core research questions: (RQ1) how feasible is WiFi CSI sensing for key tasks such as presence detection and activity recognition? To answer these questions we implement a comprehensive end-to-end pipeline encompassing five public CSI datasets UTHAR, NTU-Fi, HAR, WildActivity, and Experimental results demonstrate that 1DCNN and GRU architectures achieve competitive accuracy on standard indoor benchmarks while

"Abstract[swe]": €€€€

WiFi-avkänning utnyttjar kanalstatusinformation (CSI) från vanlig WiFi-hårdvara för att detektera och tolka fysiska aktiviteter utan behov av kameror. Trots snabba framsteg i kontrollerade laboratoriemiljöer är genomförbarheten av WiFi-CSI-avkänning under realistiska förhållanden fortfarande osäker. Denna avhandling undersöker hur genomförbart WiFi-CSI-avkänning är för närvarodetektering och grundläggande aktivitetsigenkänning under naturliga förhållanden. Experimentella resultat visar att 1DCNN och GRU uppnår konkurrenskraftig noggrannhet på standardriktmärken medan trestegspipeline o

€€€€,
"Keywords[eng]": €€€€
Wi-Fi sensing, Channel State Information, Human Activity Recognition, Deep learning, Transfer learning, Self-supervised learning, Automotive sensing, Cross-domain generalization, Convolutional Neural Network, Gated Recurrent Unit €€€€,

€€€€,
"Keywords[swe]": €€€€
Wi-Fi-avkänning, Kanalstatusinformation, Igenkänning av mänsklig aktivitet, Djupinläring, Överföringsinläring, Självövervakad inläring,
Fordonsavkänning, Korsdomänsgeneralisering €€€€,
}



acronyms.tex

```
%%% Local Variables:
%%% mode: latex
%%% TeX-master: t
%%% End:
%
% Acronyms and abbreviations for:
%   Wi-Fi CSI Sensing for Automotive In-Cabin Applications
%   KTH Royal Institute of Technology / Autoliv Research, 2026
%
% Format: \newacronym{label}{ABBREVIATION}{Full expansion}
% Entries are listed alphabetically by abbreviation.

\newacronym{ACF}{ACF}{Autocorrelation Function}

\newacronym{BVP}{BVP}{Body-Coordinate Velocity Profile}
\newacronym{BYOL}{BYOL}{Bootstrap Your Own Latent}

\newacronym{CFO}{CFO}{Carrier Frequency Offset}
\newacronym{CNN}{CNN}{Convolutional Neural Network}
\newacronym{CPD}{CPD}{Child Presence Detection}
\newacronym{CSI}{CSI}{Channel State Information}

\newacronym{DC}{DC}{Direct Current}
\newacronym{DFT}{DFT}{Discrete Fourier Transform}
\newacronym{DL}{DL}{Deep Learning}

\newacronym{ECE}{ECE}{Economic Commission for Europe}
\newacronym{EM}{EM}{Expectation--Maximisation}
\newacronym{EU}{EU}{European Union}

\newacronym{FFT}{FFT}{Fast Fourier Transform}

\newacronym{GPU}{GPU}{Graphics Processing Unit}
\newacronym{GRU}{GRU}{Gated Recurrent Unit}

\newacronym{HAR}{HAR}{Human Activity Recognition}
\newacronym{HHT}{HHT}{Hilbert--Huang Transform}
\newacronym{HMM}{HMM}{Hidden Markov Model}

\newacronym{IEEE}{IEEE}{Institute of Electrical and Electronics Engineers}
\newacronym{IoT}{IoT}{Internet of Things}
\newacronym{ISAC}{ISAC}{Integrated Sensing and Communication}

\newacronym{KTH}{KTH}{KTH Royal Institute of Technology}

\newacronym{LSTM}{LSTM}{Long Short-Term Memory}
\newacronym{LoS}{LoS}{Line of Sight}

\newacronym{ML}{ML}{Machine Learning}
\newacronym{mmWave}{mmWave}{Millimeter-Wave}
\newacronym{MoCo}{MoCo}{Momentum Contrast}

\newacronym{NIC}{NIC}{Network Interface Card}

\newacronym{OFDM}{OFDM}{Orthogonal Frequency-Division Multiplexing}

\newacronym{ReLU}{ReLU}{Rectified Linear Unit}
\newacronym{RF}{RF}{Radio Frequency}
\newacronym{RFID}{RFID}{Radio Frequency Identification}
\newacronym{RNN}{RNN}{Recurrent Neural Network}
\newacronym{RSSI}{RSSI}{Received Signal Strength Indicator}

\newacronym{SFO}{SFO}{Sampling Frequency Offset}
\newacronym{SimCLR}{SimCLR}{Simple Framework for Contrastive Learning of Representations}
\newacronym{SNR}{SNR}{Signal-to-Noise Ratio}
\newacronym{SSL}{SSL}{Self-Supervised Learning}
\newacronym{STFT}{STFT}{Short-Time Fourier Transform}

\newacronym{UN}{UN}{United Nations}
\newacronym{UNECE}{UN-ECE}{United Nations Economic Commission for Europe}

\newacronym{VICReg}{VICReg}{Variance-Invariance-Covariance Regularization}

\newacronym{WiFi}{Wi\nobreakdash-Fi}{Wireless Fidelity}
\newacronym{WLAN}{WLAN}{Wireless Local Area Network}
```